

Understanding Lvalues and Rvalues (corrected)

Dan Saks

Core C++ Conference
May, 2019

1

Lvalues and Rvalues

- *Lvalues* and *rvalues* aren't really language features.
- Rather, they're semantic properties of expressions.
- Understanding them provides valuable insights into:
 - the behavior of built-in operators
 - the code generated to execute those operators
 - the meaning of some otherwise cryptic compiler error messages
 - reference types
 - overloaded operators

5

Lvalues and Rvalues Have Evolved

- In early C, the concepts of lvalue and rvalue were fairly simple.
- Early C++ added classes, `const`, and references.
- The concepts got more complicated.
- Modern C++ added rvalue references.
- The concepts got even more complicated.
- This talk explains the origins of the concepts of lvalue and rvalue, from this historical perspective.

6

Lvalues

- In *The C Programming Language*, Kernighan and Ritchie wrote:
 - The name “lvalue” comes from the assignment expression
$$E1 = E2$$
in which the **left** operand E1 must be an lvalue expression.
 - An **lvalue** is an expression referring to an object.
 - An **object** is a region of storage.

7

Lvalues and Rvalues

```
int n;  // a definition for an integer object named n
~
n = 1;  // an assignment expression
```

- `n` is a sub-expression referring to an `int` object.
 - It's an *lvalue*.
- `1` is a sub-expression not referring to an object.
 - It's an *rvalue*.
- An *rvalue* is simply an expression that's not an lvalue.

8

Lvalues and Rvalues

- Here's a more complicated assignment:
`x[i + 1] = abs(p->value);`
- `x[i + 1]` is an expression. So is `abs(p->value)`.
- For the assignment to be valid:
 - The left operand must be an lvalue.
 - It must refer to an object.
 - The right operand can be either an lvalue or rvalue.
 - It can be any expression.

9

A Look Under the Hood

- Why make this distinction between lvalues and rvalues?
- One answer:
 - So that compilers can assume that rvalues don't necessarily occupy storage.
 - This offers considerable freedom in generating code for rvalue expressions.
- Again, let's consider the assignment in:

```
int n; // a declaration for an integer object named n
~~~
n = 1; // an assignment expression
```

10

Data Storage for Rvalues

- A compiler might represent 1 as named data storage initialized with the value 1, as if 1 were an lvalue.
- In assembly language, this might look something like:

```
one:           ; a label for the following location
    .word 1     ; allocate storage holding the value 1
```

- The compiler would generate code to copy from that initialized storage to the storage allocated for n:

```
mov n, one ; copy the value at one to location n
```

11

Data Storage for Rvalues

- Some machines provide instructions with an *immediate operand*:
 - A source operand value can be part of an instruction.
- In assembly, this might look like:

```
mov n, #1      ; copy the value 1 to location n
```
- In this case:
 - The rvalue 1 never appears as an object in the data space.
 - Rather, it appears as part of an instruction in the code space.

12

Data Storage for Rvalues

- On some machines, the preferred way to put the value 1 into an object might be to:
 - clear the object,
 - then increment it.
- In assembly, this might look like:

```
clr n      ; set n to zero  
inc n      ; increment n, effectively setting it to 1
```
- Data representing the values 0 and 1 appear nowhere in either the source or object code.

13

Must be an Lvalue == Can't be an Rvalue

- Now, suppose you write:

```
1 = n;      // obviously silly
```

- This is trying to change the value of the integer literal, 1.
- Of course, C (and C++) reject it as an error.
- But why, exactly?
 - An assignment assigns a value to an object.
 - Its left operand must be an lvalue.
 - But 1 is not an lvalue; it's an rvalue.

14

Recap

- Every expression in C is either an lvalue or an rvalue.
- In general:
 - An *lvalue* is an expression that refers to an object.
 - An *rvalue* is simply any expression that isn't an lvalue.
- Caveat: Although this is also true for non-class types in C++, it's not true for class types.

15

Literals

- Most literals are rvalues, including:
 - numeric literals, such as 3 and 3.14159
 - character literals, such as 'a'
- They don't necessarily occupy data storage.
- However, character string literals, such as "xyzy", are lvalues.
- They occupy data storage.

16

Enumeration Constants

- When used in expressions, enumeration constants are also rvalues:

```
enum color { red, green, blue };  
color c;  
~~~  
c = green;      // OK: c is an lvalue  
blue = green;   // error: blue is an rvalue
```

17

Lvalues Used as Rvalues

- An lvalue can appear on either side of an assignment, as in:

```
int m, n;  
~~~  
m = n;      // OK: m and n are both lvalues
```

- Obviously, you can assign the value in n to the object designated by m.
- This assignment uses the lvalue expression n as an rvalue.
- Officially, C++ performs an ***lvalue-to-rvalue conversion***.

18

Operands of Other Operators

- The concepts of lvalue and rvalue apply in all expressions.
 - Not just assignment.
- For example, both operands of the binary + operator must be expressions.
 - Obviously, those expressions must have suitable types.
- But each operand can be either an lvalue or rvalue.

```
int x;  
~~~  
~~~ x + 2 ~~~ // OK: lvalue + rvalue  
~~~ 2 + x ~~~ // OK: rvalue + lvalue
```

19

What About the Result?

- For built-in binary (non-assignment) operators such as `+`:
 - The operands may be lvalues or rvalues.
 - But what about the result?
- An expression such as `m + n` places its result:
 - not in `m`,
 - not in `n`,
 - but rather in a compiler-generated temporary object, often a CPU register.
- Such temporary objects are rvalues.

20

What About the Result?

- For example, this is (obviously?) an error:

```
m + 1 = n;    // error... but why?
```
- The `+` operator has higher precedence than `=`.
- Thus, the assignment expression is equivalent to:

```
(m + 1) = n;    // error... but why?
```
- It's an error because `m + 1` yields an rvalue.

21

Unary &

- `&e` is a valid expression only if `e` is an lvalue.
- Thus, `&3` is an error.
- Again, `3` does not refer to an object, so it's not addressable.
- Although the operand must be an lvalue, the result is an rvalue.
- For example,

```
int n, *p;  
~~~  
p = &n;           // OK: n is an lvalue  
&n = p;           // error: &n is an rvalue
```

22

Unary *

- In contrast to unary `&`, unary `*` yields an lvalue.
- A pointer `p` can point to an object, so `*p` is an lvalue.

```
int a[N];  
int *p = a;  
char *s = NULL;    // = nullptr in Modern C++  
~~~  
*p = 3;             // OK  
*s = '\0';          // undefined behavior
```

- Note: Lvalue-ness is a compile-time property.
 - `*s` is an lvalue even if `s` is null.
 - If `s` is null, evaluating `*s` causes undefined behavior.

23

Unary *

- Again, the result of the * operator is an lvalue.
- However, its operand can be an rvalue.
- For example,

```
*(p + 1) = 4;    // OK
```
- Here, `p + 1` is an rvalue, but `*(p + 1)` is an lvalue.
- The assignment stores the value 4 into the object referenced by `*(p + 1)`.

24

Data Storage for Expressions

- Conceptually, rvalues (of non-class type) don't occupy data storage in the object program.
 - In truth, some might.
- C and C++ insist that you program as if non-class rvalues don't occupy storage.
- Conceptually, lvalues (of any type) occupy data storage.
 - In truth, the optimizer might eliminate some of them.
 - (But only when you won't notice.)
- C and C++ let you assume that lvalues always do occupy storage.

25

Non-Modifiable Lvalues

- In fact, not all lvalues can appear on the left of an assignment.
- An lvalue is **non-modifiable** if it has a const-qualified type.
- For example,

```
char const name[] = "dan";  
~~~~  
name[0] = 'D';      // error: name[0] is const
```

- name[0] is an lvalue, but it's non-modifiable.
 - Each element of a const array is itself const.

26

Non-Modifiable Lvalues

- Lvalues and rvalues provide a vocabulary for describing subtle behavioral differences...
- ...such as between enumeration constants and const objects.
- For example, this MAX is a constant of an unnamed enumeration type:

```
enum { MAX = 100 };
```

- Unscoped enumeration values implicitly convert to integer.

27

Non-Modifiable Lvalues

- When `MAX` appears in an expression, it yields an integer rvalue.
- Thus, you can't assign to it:

```
MAX += 3;           // error: MAX is an rvalue
```

- You can't take its address, either:

```
int *p = &MAX;      // error: again, MAX is an rvalue
```

28

Non-Modifiable Lvalues

- On the other hand, this `MAX` is a const-qualified object:

```
int const MAX = 100;
```

- When it appears in an expression, it's a non-modifiable lvalue.
- Thus, you still can't assign to it.

```
MAX += 3;           // error: MAX is non-modifiable
```

- However, you can take its address:

```
int const *p = &MAX; // OK: MAX is an lvalue
```

29

Recap

- This table summarizes the behavior of lvalues and rvalues (of non-class type):

	can take the address of	can assign to
lvalue	yes	yes
non-modifiable lvalue	yes	no
(non-class) rvalue	no	no

30

Const Objects

- A const object is addressable.
 - The compiler may generate storage to hold the const object's value.
- The compiler might find that the program never needs storage for a particular const object.
 - It often does.
- In that case, the compiler need not allocate storage for that object.
- This behavior for const objects is analogous to the behavior for inline functions.

31

Reference Types

- The concepts of lvalues and rvalues help explain C++ *reference types*.
- References provide an alternative to pointers as a way of associating names with objects.
- C++ libraries often use references instead of pointers as function parameters and return types.

32

Reference Types

- Consider the following code:

```
int i;           // define i as an integer object
~~~~
int &ri = i;     // define ri as a "reference to int"
```

- The last line above:
 - defines `ri` with type “reference to `int`”, and
 - initializes `ri` to refer to `i`.
- Hence, reference `ri` is an *alias* for `i`.

33

Reference Types

- A reference is essentially a pointer that's automatically dereferenced each time it's used.
- You can rewrite most, if not all, code that uses a reference as code that uses a const pointer, as in:

reference notation	equivalent pointer notation
<code>int &ri = i;</code>	<code>int <i>*const</i> cpi = &i;</code>
<code>ri = 4;</code>	<code>*cpi = 4;</code>
<code>int j = ri + 2;</code>	<code>int j = *cpi + 2;</code>

- A reference acts like a const pointer that's dereferenced (has a * in front of it) whenever you touch it.
- A reference yields an *lvalue*.

34

Initializing vs. Assigning

- **Initializing** a reference associates the reference with an object.
 - Initializing a reference is also known as **binding**.
- **Assigning** to a reference stores through the reference and into the referenced object.
- For instance,

```
int &ri = i;    // binds reference to object
ri = 3;        // assigns to referenced object
```

35

References and Overloaded Operators

- What good are references?
- Why not just use pointers?
- References can provide friendlier function interfaces.
- More specifically, C++ has references so that overloaded operators can look just like built-in operators...

36

References and Overloaded Operators

```
enum month {  
    Jan, Feb, Mar, ~~~, Dec, month_end  
};  
typedef enum month month;  
~~~~  
  
for (month m = Jan; m <= Dec; ++m) {  
    ~~~  
}
```

- This code compiles and executes as expected in C.
- However, it doesn't compile in C++...

37

References and Overloaded Operators

- In C++, the built-in ++ won't accept an operand of enumeration type.
- You need to overload ++ for month.
- Let's try it without references...

38

References and Overloaded Operators

```
void operator++(month x) {           // pass by value
    x = static_cast<month>(x + 1);
}
```

- Using this definition, ++*m* compiles, but doesn't increment *m*.
 - It increments a copy of *m* in parameter *x*.
- Also, this implementation lets you apply ++ to an rvalue, as in:

```
++Apr;           // compiles, but shouldn't
```

- A proper overloaded ++ should behave like the built-in ++, as in:

```
++42;           // compile error: can't increment an rvalue
```

39

References and Overloaded Operators

- We need a ++ that passes in a month it can modify:

```
void operator++(month *x) {           // pass by address?
    *x = static_cast<month>(*x + 1);
}
```

- In fact, this function definition won't compile.
 - You can't overload an operator with a parameter of pointer type.
- Even if the definition compiled, it wouldn't work like a built-in ++:

```
++m;           // looks right but doesn't compile
++&m;          // looks wrong and doesn't compile
```

40

References and Overloaded Operators

- We really need a ++ that can modify a month object...
- ... but without passing explicitly by address:

```
void operator++(month &x) {           // pass by reference
    x = static_cast<month>(x + 1);
}
```

- Using this definition:

```
++m;           // compiles, increments m, and looks right
```

- As a bonus, this ++ operator won't accept an rvalue:

```
++Apr;         // compile error
```

41

References and Overloaded Operators

- Actually, a proper prefix ++ doesn't return void.
- It returns the incremented object by reference:

```
month &operator++(month &x) {           // non-void return  
    return x = static_cast<month>(x + 1);  
}
```

- This enables overloaded ++ to act even more like a built-in operator:

```
int j = ++i;    // OK  
month n = ++m; // OK
```

42

“Reference to Const” Parameters

- Just as you can have “pointer to const” parameters...
- You can also have “reference to const” parameters:

```
R f(T const &t);
```

- A “reference to const” parameter will accept an argument that's either const or non-const.
- In contrast, a reference (to non-const) parameter will accept only a non-const argument.
- When it appears in an expression, a “reference to const” yields a *non-modifiable lvalue*.

43

“Reference to Const” Parameters

- For the most part, a function declared as:

```
R f(T const &t);    // by "reference to const"
```

has the same outward behavior as a function declared as:

```
R f(T t);          // by value
```

- That is, the calls look and act very much the same...

44

“Reference to Const” Parameters

- Either way you declare `f`, you write the argument expression the same way:

```
T x;
```

~~~~

```
f(x);    // by value, or by "reference to const"?
```

- Either way, calling `f` can't alter the actual argument, `x`:
  - By value: `f` has access only to a copy of `x`, not `x` itself.
  - By “reference to const”: `f`'s parameter is declared to be non-modifiable.

45

## Why Use “Reference to Const”?

- Why pass by “reference to const” instead of by value?
- Passing by “reference to const” might be much more efficient than passing by value.
- It depends on the cost to make a copy.

46

## References and Temporaries

- A “pointer to T” can point only to an lvalue of type T.
- Similarly, a “reference to T” binds only to an lvalue of type T.
- For example, these are both compile errors:

```
int *pi = &3;      // can't apply & to 3
int &ri = 3;       // can't bind this, either
```

- These are also compile errors:

```
int i;
~~~
double *pd = &i; // can't convert pointers
double &rd = i; // can't bind this, either
```

47

## References and Temporaries

- There's an exception to the rule that a reference must bind to an lvalue of the referenced type:
  - A "reference to const T" can bind to an expression x that's not an lvalue of type T ...
  - ... if there's a conversion from x's type to T.
- In this case, the compiler creates a temporary object to hold a copy of x converted to T.
  - This is so the reference has something to bind to.

48

## References and Temporaries

- For example, consider:  

```
int const &ri = 3;
```
- When program execution reaches this declaration, the program:
  1. creates a temporary `int` to hold the 3, and
  2. binds `ri` to the temporary.
- When execution leaves the scope containing `ri`, the program:
  3. destroys the temporary.

49

## References and Temporaries

- Given:

```
double const &rd = ri; // ri from the previous slide
```

- When program execution reaches this declaration, the program:
  1. converts the value of `ri` from `int` to `double`,
  2. creates a temporary `double` to hold the converted result, and
  3. binds `rd` to the temporary.
- Again, when execution leaves the scope containing `rd`, the program:
  4. destroys the temporary.

50

## References and Temporaries

- This special behavior enables passing by “reference to const” to consistently have the same outward behavior as passing by value.
- For example, compare this with the code on the next slide:

```
long double x;
void f(long double ld); // by value
~~~
```

```
f(x); // passes a copy of x  
f(1); // passes a copy of...  
      // ...1 converted to long double
```

51



## References and Temporaries

- This is the same example, except it uses a “reference to const” parameter in place of a value parameter:

```
long double x;  
void f(long double const &ld); // by reference to const  
~~~  

f(x); // passes a reference to x
f(1); // passes a reference to a temporary...
 // ...containing 1 converted to long double
```

- Either way, the function calls behave the same.

52

## Mimicking Built-In Operators

- Recall the behavior of the built-in + operator:
  - The operands may be lvalues or rvalues.
  - The result is always an rvalue.
- How do you declare an overloaded operator with the same behavior?
- Consider a rudimentary (character) string class with + as a concatenation operator...

53

## Mimicking Built-In Operators

- You can declare operator + as a non-member, as in:

```
class string {
public:
 string(string const &);
 string(char const *); // converting constructor
 string &operator=(string const &);
    ~~~
};

string operator+(string const &lo, string const &ro);
```

54

## Mimicking Built-In Operators

```
string operator+(string const &lo, string const &ro);
```

- Parameters lo and ro accept arguments that are either lvalues or rvalues:

```
string s = "hello";
string t = "world";
~~~
s = s + ", " + t;
```

- The compiler applies the converting constructor implicitly:

```
s = s + string(", ") + t; // lvalue + rvalue + lvalue
```

55

## Mimicking Built-In Operators

```
string operator+(string const &lo, string const &ro);
```

- The function returns its result by value.
- Calling this operator + yields an rvalue:

```
string *p = &(s + t); // error: can't take the address
```

56

## References

- C++11 introduced another kind of reference.
- What C++03 calls “*references*”, C++11 calls “*lvalue references*”.
- This distinguishes them from C++11’s new “*rvalue references*”.
- Except for the name change, lvalue references in C++11 behave just like references in C++03.

57

## Rvalue References

- Whereas an *lvalue reference* declaration uses the `&` operator, an *rvalue reference* uses the `&&` operator.
- For example, this declares `ri` to be an “rvalue reference to `int`”:

```
int &&ri = 10;
```

- You can use rvalue references as function parameters and return types, as in:

```
double &&f(int &&ri);
```

- You can also have an “rvalue reference to `const`”, as in:

```
int const &&rci = 20;
```

58

## Rvalue References

- Rvalue references bind only to rvalues.
- This is true even for “rvalue reference to `const`”.
- For example,

```
int n = 10;
int &&ri = n; // error: n is an lvalue
int const &&rj = n; // error: n is an lvalue
```

59

## Move Operations

- Modern C++ uses rvalue references to implement move operations that can avoid unnecessary copying:


```
class string {
public:
 // copy operations
 string(string const &); // constructor
 string &operator=(string const &); // assignment

 // move operations
 string(string &&) noexcept; // constructor
 string &operator=(string &&) noexcept; // assignment
};
```

60

## Value Categories

- Modern C++ introduces a more complex categorization of expressions:



- The newer categories are:
  - glvalue**: a “generalized” lvalue
  - prvalue**: a “pure” rvalue
  - xvalue**: an “expiring” lvalue

61

Thanks for Listening